

Advanced Computer Networks

Module III

Syllabus

- Transport Layer and Network Layer Security: Connection oriented and connection-less services, Transmission Control Protocol and User Datagram Protocol – characteristics, reliability mechanism, performance considerations. Flow Control - Go-back N and Selective Repeat sliding window protocols and Congestion control mechanisms. IP Addressing - IPv4 Classful addressing, Subnet masking and Packet format. Unicast routing algorithms – Link- State and Distance Vector routing, Routing protocols – RIP and OSPF. SSL for secure web communication – handshake process and vulnerabilities, IPSec for secure IP communication and firewalls – types, Network Address Translation management.

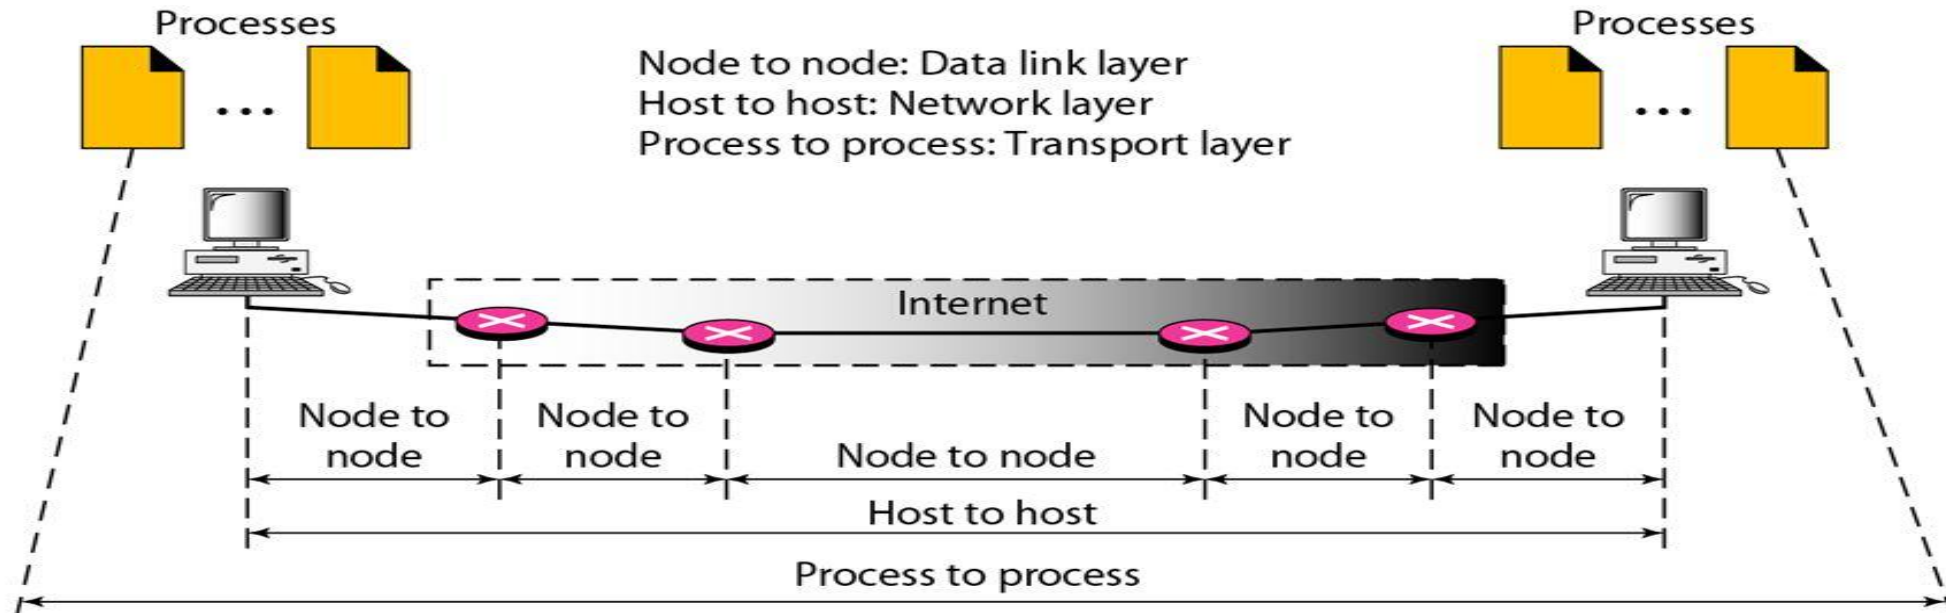
Transport Layer and Network Layer Security - Introduction

- Network Layer Responsibility: Ensures delivery of datagrams between two hosts (host-to-host delivery).
- Internet communication is not just between nodes or hosts but between processes (application programs).
- So there is a need for Process-to-Process delivery.
- Real communication happens between processes, not just hosts.
- Several processes may run on both the source and destination hosts simultaneously.
- A mechanism is required to ensure data is delivered from the correct source process to the corresponding destination process.

Continues..

- Transport Layer ensures process-to-process delivery of packets.
- It facilitates data transfer between processes rather than just hosts.
- Processes communicate in a client-server relationship.
- There are three types of data deliveries:- Node to Node, Host to Host and Process to Process

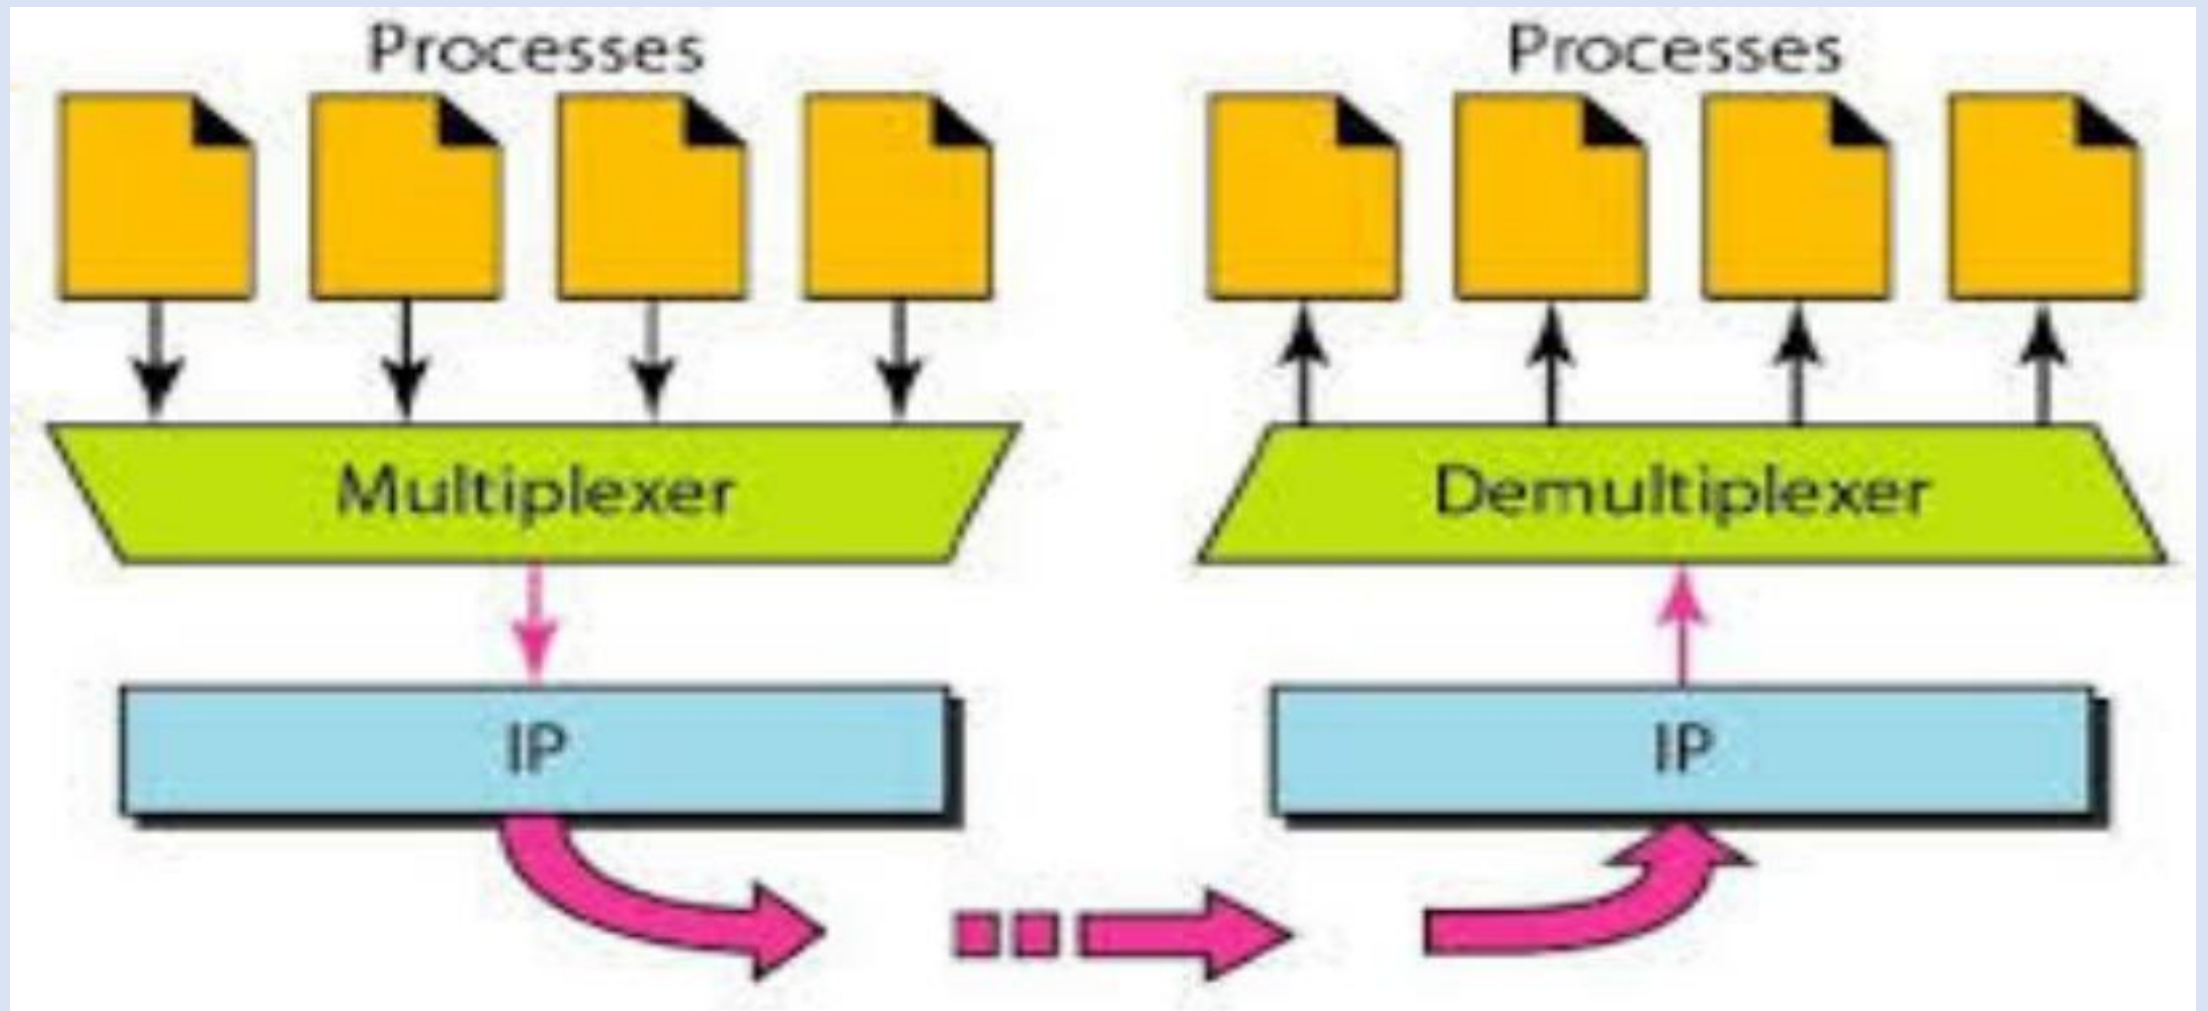
Types of data deliveries



Multiplexing and Demultiplexing

- Multiplexing (Sender Side)
 - Multiple processes need to send packets.
 - Only one transport layer protocol is available at a time.
 - The many-to-one relationship requires multiplexing.
 - Messages from different processes are identified by their port numbers.
 - The transport layer adds a header and forwards the packet to the network layer.
- Demultiplexing (Receiver Side)
 - The transport layer receives datagrams from the network layer.
The one-to-many relationship requires demultiplexing.
 - Error checking is performed, and the header is removed.
 - The message is delivered to the appropriate process based on the port number.

Continues..



Connection-Oriented and Connectionless services

Difference Between Connection-Oriented and Connectionless Services

| Feature | Connection-Oriented Service | Connectionless Service |
|------------------|-----------------------------------|---------------------------------|
| Connection Setup | Required before data transfer | Not required |
| Packet Ordering | Ensures ordered delivery | May arrive out of order |
| Reliability | Reliable (uses acknowledgments) | Unreliable (no acknowledgments) |
| Error Handling | Error checking and retransmission | No error correction |
| Speed | Slower due to overhead | Faster due to minimal control |
| Examples | TCP, SCTP | UDP |

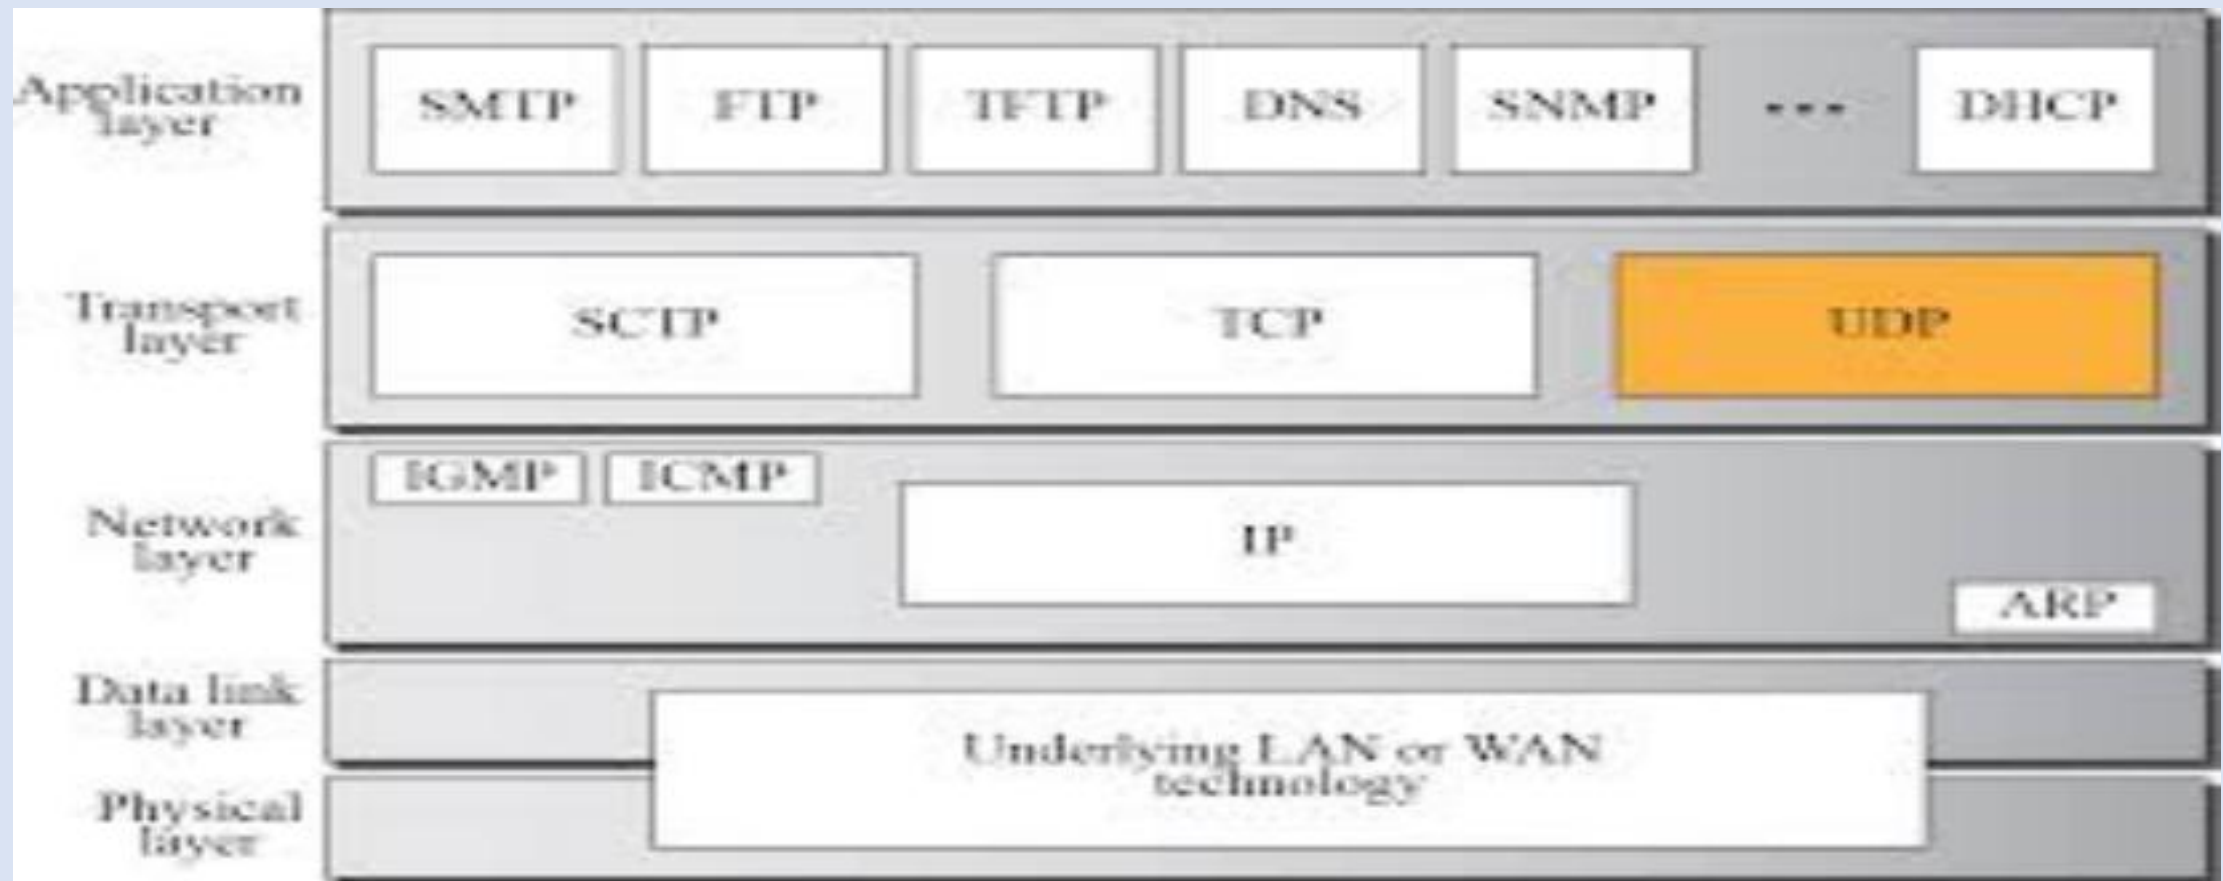
Continues..

- Reliable vs. Unreliable Transport Layer
- The transport layer can provide reliable or unreliable services based on application needs.
- Reliable Transport Service:
 - Implements flow and error control to ensure accurate data delivery.
 - More complex and slower due to additional processing.
 - Used when applications require guaranteed data integrity.
- Unreliable Transport Service:
 - No built-in flow or error control.
 - Faster but may result in packet loss or errors.
 - Suitable for real-time applications or when the application manages reliability.

Continues..

- Transport Layer Protocols in the Internet:
 - UDP: Connectionless and unreliable.
 - TCP & SCTP: Connection-oriented and reliable.
- Why Reliability at the Transport Layer?
 - The Data Link Layer ensures reliability only between two nodes.
 - The Network Layer provides best-effort delivery (unreliable).
 - Transport Layer reliability ensures end-to-end error-free communication.

Position of UDP, TCP, and SCTP in TCP IP suite



User Datagram Protocol

- Connectionless unreliable transport protocol.
- Provides process to process communication.
- Simple protocol.
- Use minimum overhead.
- Message passing with less reliability can use UDP protocol.
- Why UDP ?
 - Finer application-level control over what data is sent, and when.
 - No connection establishment.
 - No connection state.
 - Small packet header overhead

Continues..

- UDP takes messages from the application process, attaches source and destination port number fields for the multiplexing / demultiplexing service, adds two other small fields, and passes the resulting segment to the network layer.
- The network layer encapsulates the transport-layer segment into an IP datagram and then makes a best-effort attempt to deliver the segment to the receiving host.
- If the segment arrives at the receiving host, UDP uses the destination port number to deliver the segment's data to the correct application process.
- With UDP there is no handshaking between sending and receiving transport-layer entities before sending a segment.
- For this reason, UDP is said to be connectionless.

Well-known ports used with UDP

- Some port numbers can be used by both UDP and TCP.
- Well-known ports are predefined port numbers (0-1023) assigned by IANA (Internet Assigned Numbers Authority).
- They are used by common network services and protocols like HTTP (port 80), DNS (port 53), and DHCP (port 67/68).
- These ports help in identifying and standardizing communication between applications across networks.
- Why well-known ports?
 - Standardized communication across networks.
 - Service identification for specific applications.
 - Interoperability between systems and devices.
 - Efficient routing and traffic management.
 - Security enforcement through firewalls and access control.
 - Simplifies application development by avoiding conflicts.
 - Quick troubleshooting for network administrators.

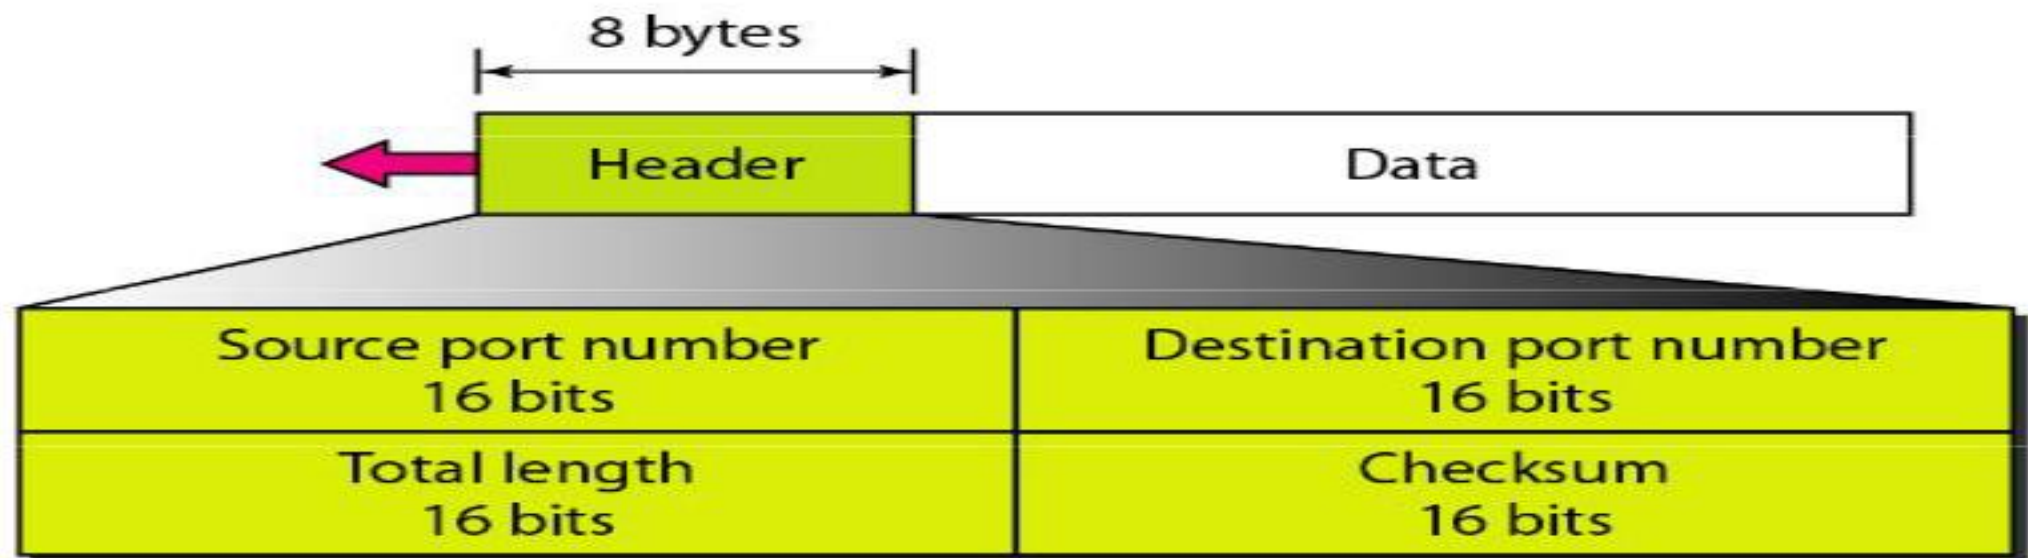
Well-known ports used with UDP

Table 4.1 Well-known port numbers used by UDP

| Port | Protocol | Description |
|------|-------------|---|
| 1 | Echo | Echoes a received datagram back to the sender |
| 9 | Discard | Discards any datagram that is received |
| 11 | Users | Active users |
| 13 | Daytime | Returns the date and the time |
| 17 | Quote | Returns a quote of the day |
| 19 | Chargen | Returns a string of characters |
| 53 | Name server | Domain Name Service |
| 67 | BOOTPs | Server port to download bootstrap information |
| 68 | BOOTPc | Client port to download bootstrap information |
| 69 | TFTP | Trivial File Transfer Protocol |
| 111 | RPC | Remote Procedure Call |
| 123 | NTP | Network Time Protocol |
| 161 | SNMP | Simple Network Management Protocol |
| 162 | SNMP | Simple Network Management Protocol (trap) |

Continues..

User datagram format

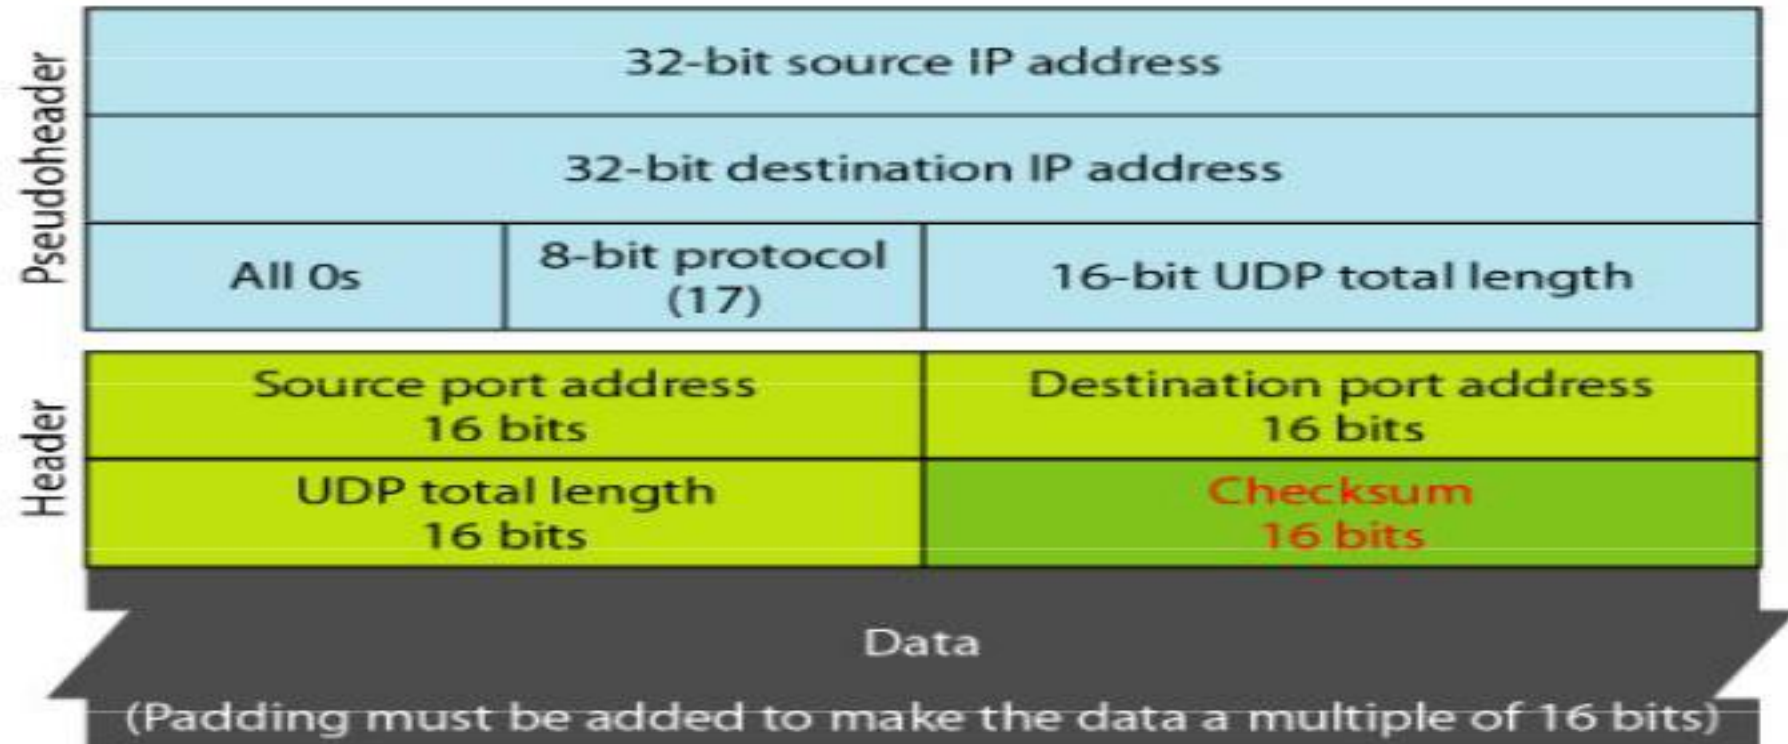


Continues..

- User Datagram
 - UDP packets are called user datagram.
 - Fixed size header(8 bytes)
 - Four fields (2 bytes)
 - First two fields : source and destination port address
 - Third field : total length of the user datagram, header plus data. The 16 bits can define a total length of 0 to 65535 bytes.
 - The last bit may carry optional checksum.

Continues..

Pseudoheader for checksum calculation



Continues..

- Process to process communication.
- Using socket address: a combination of IP address and port address.
- Connectionless Service.
- Each datagram send by UDP is an independent datagram.
- Packets from same source process to same destination process are also independent to each other.
- The user datagram are not numbered.
- No connection establishment and termination.
- User datagram can take different path.
- Disadvantages of UDP protocol.
- Short messages less than 65,507 bytes can use UDP.

Continues..

- Flow Control
 - Simple protocol.
 - No flow control.
 - The receiver may overflow with incoming messages.
 - The process using this UDP should provide this service.
- Error Control
 - No error control mechanism except for checksum.
 - The sender does not know if a message is lost or is duplicated.
 - When an error is detected through the checksum, the user datagram is silently discarded.
 - The process using this UDP service should provide this.

Continues..

- Checksum
 - The UDP checksum provides for error detection.
 - The checksum is used to determine whether bits within the UDP segment have been altered as it moved from source to destination.
 - UDP at the sender side performs the 1s complement of the sum of all the 16-bit words in the segment, with any overflow encountered during the sum being wrapped around. This result is put in the checksum field of the UDP segment.

Continues..

- Uses of UDP
 - Ideal for simple request-response communication with minimal error control.
 - Suitable for applications with internal flow and error control, like TFTP.
 - Supports multicasting, which TCP does not.
 - Used in network management (e.g., SNMP). Helps in routing updates, such as RIP.

Transmission Control Protocol

- TCP (Transmission Control Protocol) is a core transport layer protocol in the Internet model.
- Process-to-process communication.
- It ensures reliable, error-free, and in-order data delivery.
- Unlike UDP, TCP establishes a connection before transmitting data.
- Example: When you load a webpage, your browser establishes a TCP connection with the web server to ensure all elements (HTML, images, scripts) are received correctly.

Continues..

- Why TCP is Important?
 - Ensures Reliable Communication: Every packet is acknowledged to prevent data loss.
 - Manages Network Congestion: Dynamically adjusts speed to avoid overloading the network.
 - Guarantees Ordered Delivery: Packets arrive in sequence, avoiding jumbled data.
 - Example: When downloading a file, TCP ensures each chunk arrives in order and is complete.

Continues..

- TCP Services


- Reliable Delivery: Resends lost packets, ensuring complete data transmission.
- Connection-Oriented Communication: Establishes a session before sending data.
- Flow Control: Regulates data flow to match the receiver's processing speed.
- Error Control: Uses checksums and acknowledgments (ACKs) to detect and correct errors.
- Example: Sending an email via SMTP (which runs on TCP) ensures every part of the message is delivered accurately.

TCP Well Known Ports

| <i>Port</i> | <i>Protocol</i> | <i>Description</i> |
|-------------|-----------------|---|
| 7 | Echo | Echoes a received datagram back to the sender |
| 9 | Discard | Discards any datagram that is received |
| 11 | Users | Active users |
| 13 | Daytime | Returns the date and the time |
| 17 | Quote | Returns a quote of the day |
| 19 | Chargen | Returns a string of characters |
| 20 | FIP, Data | File Transfer Protocol (data connection) |
| 21 | FIP, Control | File Transfer Protocol (control connection) |
| 23 | TELNET | Tenninal Network |
| 25 | SMTP | Simple Mail Transfer Protocol |
| 53 | DNS | Domain Name Server |
| 67 | BOOTP | Bootstrap Protocol |
| 79 | Finger | Finger |
| 80 | HTTP | Hypertext Transfer Protocol |
| 111 | RPC | Remote Procedure Call |

Continues..

- TCP treats data as a continuous stream of bytes, unlike UDP's discrete messages.
- TCP provides a virtual connection between sender and receiver, simulating a "tube" for seamless data flow.
- The sending process writes data to the stream, while the receiving process reads from it continuously.



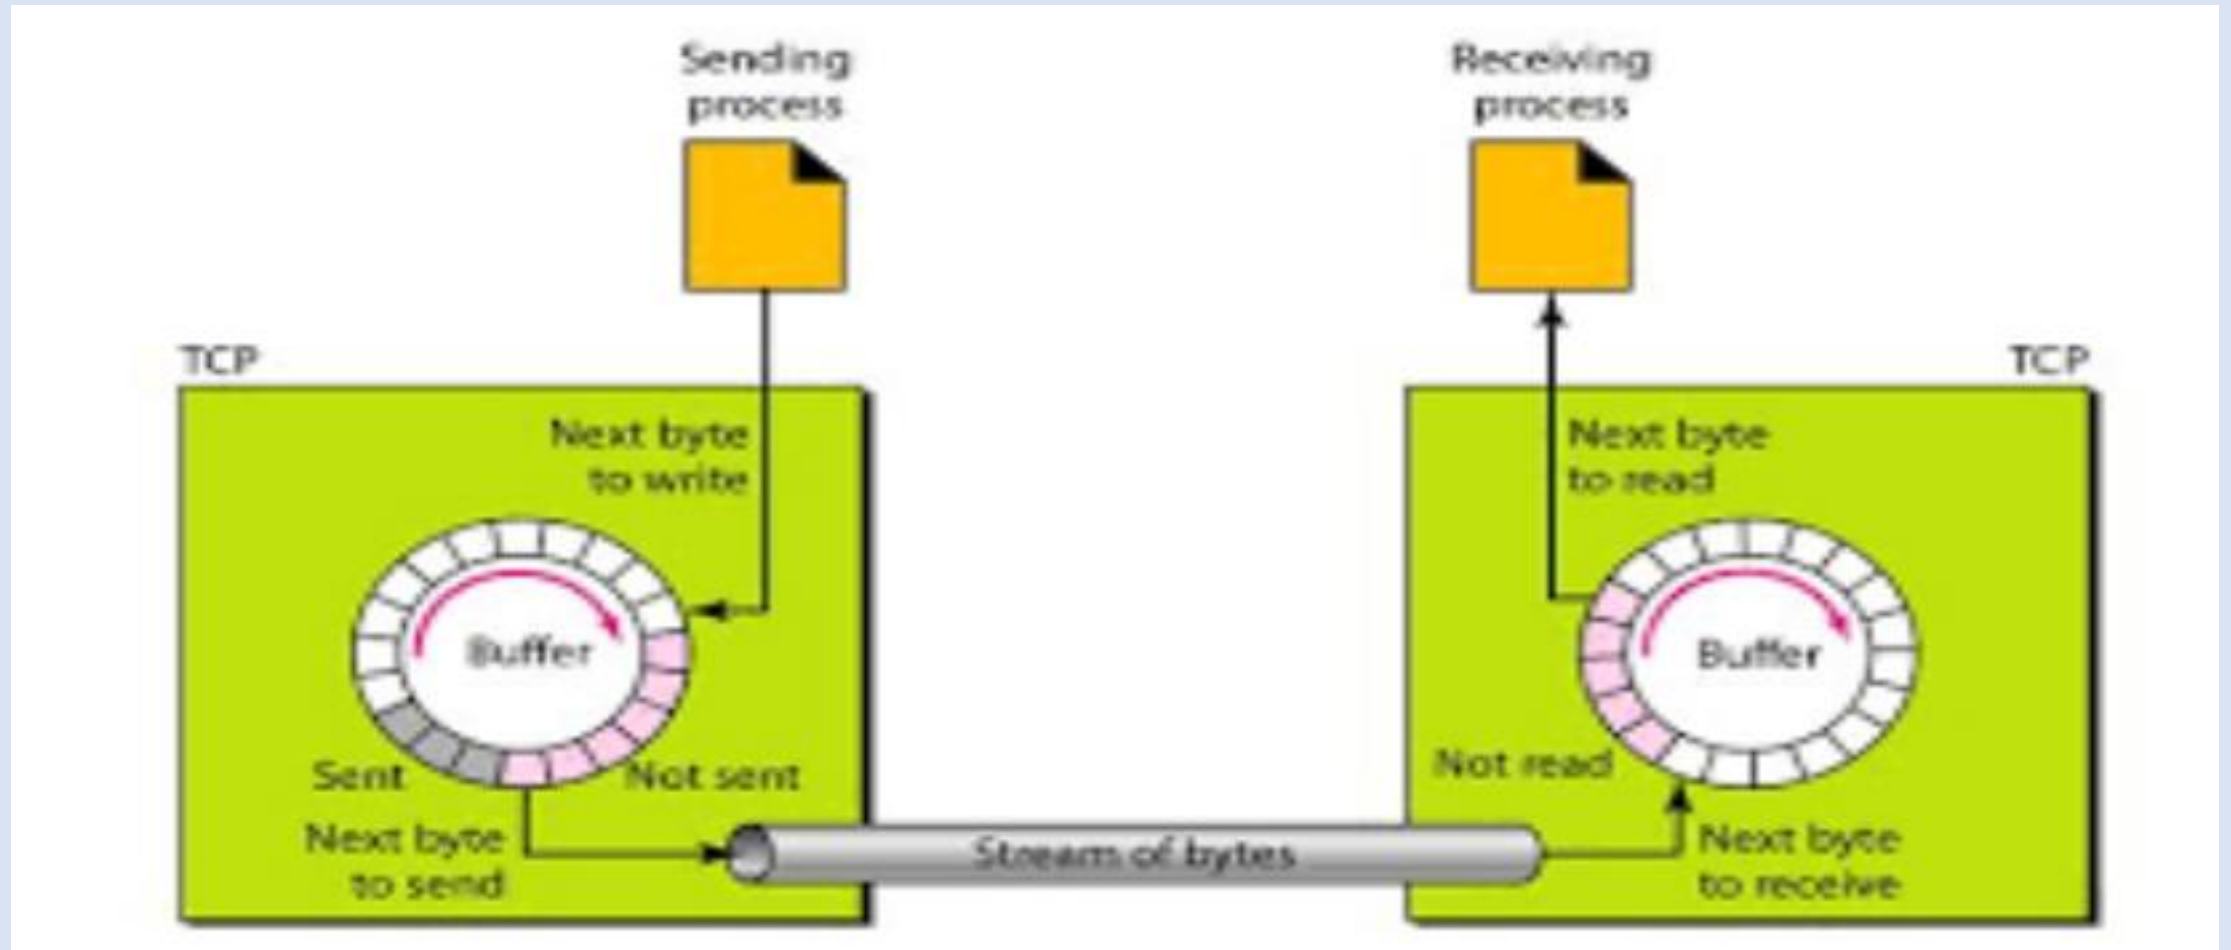
Continues..

- TCP uses buffers to handle differences in the speed of data writing (sending) and reading (receiving).
- Each direction has its own buffer—one for sending and one for receiving.
- Buffers help with flow and error control in TCP communication.
- A circular array is one way to implement these buffers, storing data in a continuous loop.
- Buffer sizes vary—they can be hundreds or thousands of bytes and are not always equal in size.

Continues..

- Sending buffer has three sections:
 - Empty chambers (ready for new data).
 - Sent but unacknowledged bytes (held until acknowledged).
 - Bytes ready to be sent (may be sent partially due to network conditions).
- Acknowledged bytes are recycled, making space for new data in the circular buffer.
- Receiving buffer has two sections:
 - Empty chambers (for incoming data).
 - Received bytes (ready to be read by the receiving process).
- TCP segments data before sending, grouping bytes into packets called segments.
- Segments are encapsulated in IP datagrams for transmission, and TCP handles issues like reordering, loss, or corruption transparently.

Continues..



Full-Duplex Communication

- TCP supports full-duplex communication, allowing data to flow simultaneously in both directions.
- Each TCP connection has both a sending and receiving buffer, enabling bidirectional data transfer.
- Segments move in both directions independently, ensuring continuous and efficient communication.

Continues..

- Connection-oriented service
 - TCP establishes a connection before data exchange and terminates it afterward.
 - The connection is virtual, not physical, as data travels in IP datagrams that may take different paths.
 - TCP ensures ordered and reliable delivery, even if packets are lost, out of order, or corrupted.
 - TCP creates a stream-oriented environment, similar to a bridge connecting two sites for seamless communication.
- Reliable service
 - TCP ensures reliable data transmission.
 - Acknowledgment Mechanism: It verifies the successful arrival of data.
 - Error Control

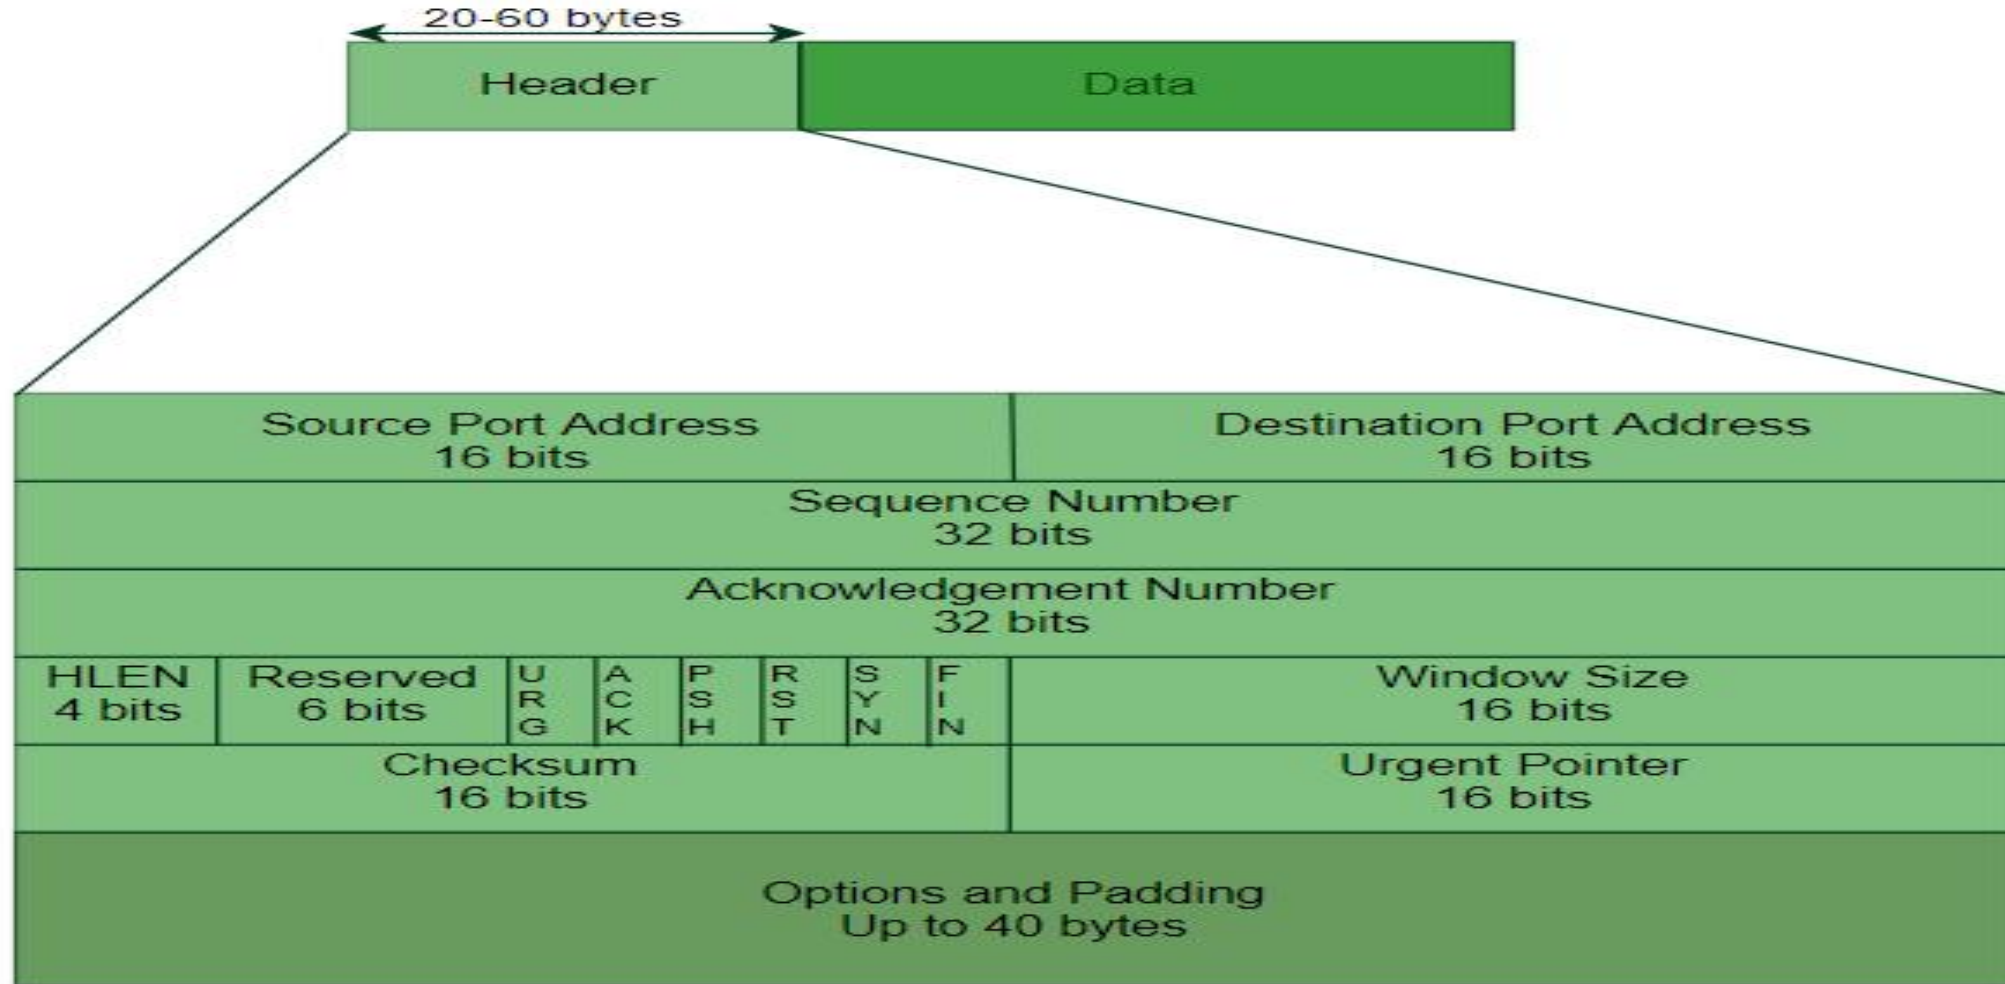
TCP Features

- No Segment Number Field: TCP tracks transmitted/received segments but does not use a segment number field in the header.
- Sequence & Acknowledgment Numbers: These fields represent byte numbers, not segment numbers.
- Byte Numbering: TCP numbers all data bytes independently in each direction.
- Random Starting Number: The first byte number is randomly chosen between 0 and $2^{32}-1$.
- Flow & Error Control: Byte numbering is crucial for managing flow and error control.
- Sequence Number: Each segment's sequence number is the number of its first byte.
- Example
 - Random Start: The first byte number is randomly chosen (e.g., 1057).
 - Byte Range: If 6000 bytes are sent, they are numbered from 1057 to 7056.

Features of TCP

- Flow Control: TCP prevents the receiver from being overwhelmed by regulating the data sent.
- Byte-Oriented Control: TCP uses a numbering system for byte-based flow control.
- Error Control: Ensures reliable transmission by detecting lost or corrupted segments.
- Segment-Based Detection: Errors are detected at the segment level but managed at the byte level.
- Congestion Control: TCP adjusts data transmission based on network congestion, unlike UDP.

Segment Format



Continues..

- Segment Structure: Consists of a header (20–60 bytes) and application data.
- Header Size: 20 bytes without options, up to 60 bytes with options.
- Fields
 - Source Port Address: 16-bit field specifying the sender's application port.
 - Destination Port Address: 16-bit field specifying the receiver's application port.
 - Sequence Number:
 - 32-bit field indicating the first byte number in the segment.
 - Byte Numbering: TCP numbers each byte to ensure reliable transmission.
 - Initial Sequence Number (ISN): Randomly generated during connection setup, different for each direction.

Continues..

- Acknowledgment Number Field:
 - A 32-bit field in the TCP header.
 - Specifies the next expected byte from the sender.
 - How It Works:
 - If the receiver successfully receives byte x , it sends $x + 1$ as the acknowledgment number.
 - Ensures reliable communication and prevents data loss or duplication.
 - Example Scenario:
 - Sender → Receiver: Sends 1000 bytes starting from byte 5000 (bytes 5000 to 5999).
 - Receiver → Sender: Acknowledges the reception by sending ACK = 6000.
 - Sender → Receiver: Sends bytes 6000–6999.
 - Receiver → Sender: Sends ACK = 7000 to indicate the next expected byte.

Continues..

- Header Length:
 - A 4-bit field indicating the number of 4-byte words in the TCP header.
 - The header length ranges from 20 to 60 bytes.
 - Field value can be between 5 ($5 \times 4 = 20$ bytes) and 15 ($15 \times 4 = 60$ bytes).
- Reserved:
 - A 6-bit field reserved for future use.
 - Currently not used but may be allocated in future versions of TCP.
- Control:
 - A field that defines 6 control bits/flags.
 - These flags control various functions in TCP communication (e.g., SYN, ACK, FIN).

Continues..

Table 4.3 Description of flags in the control field

| Flag | Description |
|------|---|
| URG | The value of the urgent pointer field is valid. |
| ACK | The value of the acknowledge field is valid. |
| PSH | Push the data. |
| RST | Reset the connection. |
| SYN | Synchronize sequence numbers during connection |
| FIN | Terminate the connection. |

Continues..

- Window Size:
 - The window size in TCP is used for flow control to manage the amount of data that can be sent before receiving an acknowledgment.
 - It helps in preventing congestion and ensures that the receiver's buffer is not overwhelmed by too much data.
 - Defines the size of the window (in bytes) that the receiver maintains.
 - The field is 16 bits long, allowing a maximum window size of 65,535 bytes.
 - Referred to as the receiving window (rwnd), determined by the receiver.
 - The sender must comply with the receiver's window size.

Continues..

- Checksum:
 - A 16-bit field used for error-checking.
 - The checksum is mandatory in TCP (unlike in UDP, where it's optional).
 - The same checksum calculation method as UDP is used, including a pseudoheader with a protocol field value of 6.
- Urgent Pointer:
 - A 16-bit field, valid only if the urgent flag is set.
 - Indicates the last byte of urgent data by adding the value to the sequence number.
- Options:
 - Up to 40 bytes of optional information can be included in the TCP header.


TCP Connection

- TCP as a Connection-Oriented Protocol
 - Establishes a virtual path between source and destination.
 - All message segments travel through this virtual path.
 - Ensures acknowledgment and retransmission of lost or damaged segments.
- Three Phases of a TCP Connection
 - Connection Establishment
 - TCP operates in full-duplex mode.
 - Both parties must agree before data transfer begins.
 - Data Transfer
 - Ensures reliable, ordered delivery of segments.
 - Retransmits lost or corrupted segments.
 - Connection Termination
 - Properly closes the communication channel after data transfer.

Continues..

- Three-Way Handshaking (Connection Establishment Process)
 - Step 1: Server initiates a passive open, indicating readiness to accept connections.
 - Step 2: Client requests an active open, signaling the need for a connection.
 - Step 3: TCP initiates a three-way handshake to establish the connection.

Continues..



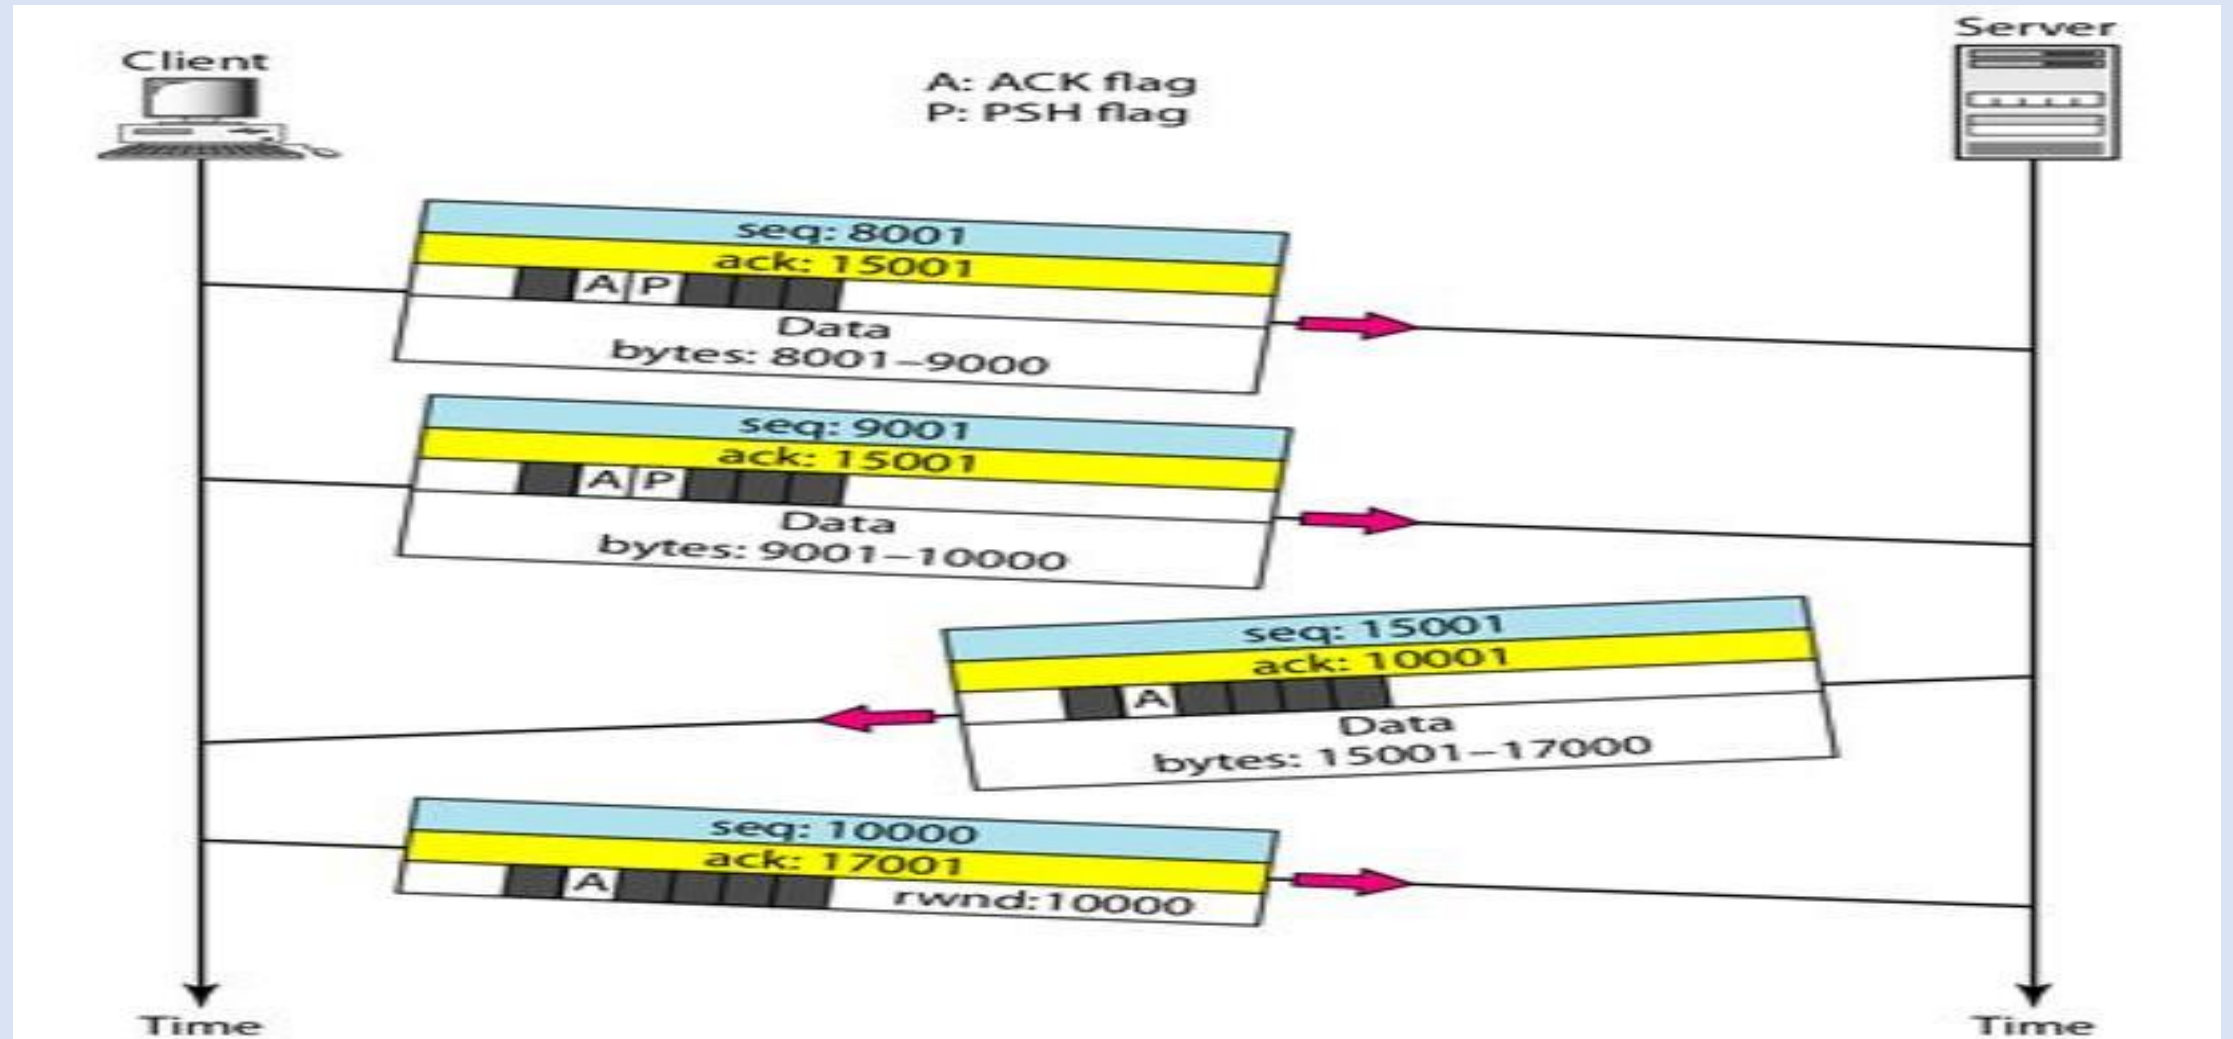
Continues..

- Step 1:
 - The client sends a SYN segment with only the SYN flag set.
 - Used to synchronize sequence numbers with the server.
 - Consumes one sequence number, incrementing it by 1 for data transfer.
 - Though it carries no real data, it is treated as having one imaginary byte.
- Step 2:
 - The server sends a SYN + ACK segment with both SYN and ACK flags set.
 - This segment acknowledges the client's SYN request.
 - It also acts as a SYN segment to initiate communication in the other direction.
 - The segment consumes one sequence number.

Continues..

- Step 3:
 - The client sends the third segment, which is an ACK segment.
 - It acknowledges the second segment using the ACK flag and acknowledgment number.
 - The sequence number remains the same as in the SYN segment.
 - The ACK segment does not consume any sequence numbers.
- **What is Simultaneous Open?**
 - Occurs when both processes issue an **active open** simultaneously.
 - Both sides send a **SYN + ACK** segment to each other.
 - A **single connection** is established between them.

Data Transfer



Continues..

- Bidirectional Data Transfer
 - Once the connection is established, both client and server can send data and acknowledgments.
 - Acknowledgments can be piggybacked with data traveling in the same direction.
 - The PSH (push) flag is used to indicate immediate data delivery.


Continues..

- Example Scenario Explanation
 - Client Sends Data (2000 bytes in Two Segments)
 - The client has 2000 bytes of data to send.
 - It divides this data into two segments (e.g., 1000 bytes each).
 - Each segment contains both data and an acknowledgment for previous server messages (if any).
 - Server Responds with 2000 Bytes in One Segment
 - The server receives both segments from the client.
 - It then sends 2000 bytes of its own data in one single segment instead of splitting it.
 - This segment also includes an acknowledgment for the client's data.
 - Client Sends a Final Acknowledgment
 - After receiving the server's 2000-byte segment, the client has no more data to send.
 - It still needs to acknowledge the received data.
 - It sends a final segment that contains only an acknowledgment (ACK) with no data.

Continues..

- Urgent Data
 - TCP is stream-oriented, treating data as a continuous stream of bytes.
 - Each byte has a position in the stream.
 - Sometimes, an application needs to send urgent data that must be processed out of order.
 - Urgent data allows the receiving application to prioritize specific bytes.
 - For that send a segment with the URG bit set.

Connection Termination



Continues..

- Connection Termination can be initiated by either the client or server, but it's usually initiated by the client.

Three-Way Handshaking for Connection Termination

- Step 1
 - The **client TCP** sends a **FIN segment** after receiving a close command from the client process.
 - The **FIN segment** may include the **last chunk of data** from the client.
 - This marks the start of the termination process in a normal situation.

Continues..

- Step 2

- After receiving the FIN segment, the server TCP informs its process.
- The server sends a FIN + ACK segment to:
 - Confirm receipt of the FIN segment from the client.
 - Announce the closure of the connection in the other direction.
- The FIN + ACK segment may also include the last chunk of data from the server.
- If no data is included, the segment consumes only one sequence number.

- Step 3

- The client TCP sends the last segment, an ACK segment.
- This segment confirms receipt of the FIN segment from the server.
- The acknowledgment number is 1 plus the sequence number received in the server's FIN segment.
- The ACK segment cannot carry data and consumes no sequence numbers.

Flow Control

- What is Flow Control?
 - Flow control is a technique used to manage the rate of data transmission between two devices in a network.
 - It ensures the sender doesn't overload the receiver with too much data.
- Why is Flow Control Important?
 - Prevents buffer overflow at the receiving end.
 - Ensures efficient and reliable communication between devices.
 - Helps avoid data loss and retransmissions.
- Types of Flow Control Mechanisms:
 - Stop-and-Wait: The sender waits for an acknowledgment after sending each frame.
 - Sliding Window: Allows the sender to send multiple frames before needing an acknowledgment.
 - Credit-Based Flow Control: The receiver grants permission (credits) to the sender on how many packets it can send.

Continues..

- Flow Control in TCP:
 - Uses a sliding window to control the flow of data.
 - The window size adjusts dynamically based on the receiver's buffer capacity.
- Key Benefits of Flow Control:
 - Prevents data congestion.
 - Enhances network reliability.
 - Optimizes bandwidth utilization.

Sliding Window Protocol