

ADVANCED OPERATING SYSTEMS

MINU MARIA JOSEPH

Module 5 [7 Hours]

Database Systems: Requirements of a Database Operating System, Problem of Concurrency Control, Serializability, Basic Synchronization Primitives for Concurrency Control- Lock Based Algorithms-Static Locking, Two-Phase Locking (2PL), Time Stamp Based Algorithms- Basic Timestamp Ordering Algorithm, Thomas Write Rule (TWR), Multiversion Timestamp Ordering Algorithm, Conservative Timestamp Ordering Algorithm, Optimistic Algorithms.

Self-Study: Computer security and database security.

-
- A concurrency control algorithm controls the interleaving of conflicting actions of transactions so that the integrity of a database is maintained, i.e., their net effect is a serial execution.

Locks

- In lock based techniques, each data object has a lock associated with it
- A transaction can request, hold, or release the lock on a data object.
- When a transaction holds a lock, the transaction is said to have locked the corresponding data object.
- A transaction can lock a data object in two modes: exclusive and shared.
- If a transaction has locked a data object in exclusive mode, no other transaction can concurrently lock it in any mode.
- If a transaction has locked a data object in shared mode, other transactions can concurrently lock it but only in shared mode.
- Basically, by locking data objects, a transaction ensures that the locked data objects are inaccessible to other transactions, while temporarily in inconsistent states.

Timestamp

- A timestamp is a unique number that is assigned to a transaction or a data object and is chosen from a monotonically increasing sequence.
- Timestamps are commonly generated according to Lamport's scheme .
- Every site S_i has a logical clock C_i , which takes integer values.
- When a transaction T is submitted at a site S_i ,
- S_i increments C_i by one and then assigns a 2-tuple (C_i, i) to T .
- The 2-tuple is referred to as the timestamp of T and is denoted by $TS(T)$.

Timestamp

- Every message contains the current clock value of its sender site,
- When a site S_j receives a message with clock value t , it sets C_j to $\max(t + 1, C_j)$.
- For any two timestamps
 - ✓ $ts_1 = (t_1, i_1)$ and
 - ✓ $ts_2 = (t_2, i_2)$,
 - ✓ $ts_1 < ts_2$, if either $(t_1 < t_2)$, or $(t_1 = t_2 \text{ and } i_1 < i_2)$

Timestamp

➤ Timestamps have two properties:

✓ uniqueness

- Timestamps generated by different sites differ in their site id part and
- Timestamps generated by the same site differ in their clock value part and

Timestamp

- ✓ Monotonicity
 - a site generates timestamps in increasing order.
 - Timestamps allow us to place a total ordering on the transactions of a distributed database system by simply ordering the transactions by their timestamps.
 - In concurrency control algorithms for distributed database systems, whenever two concurrent transactions conflict, all sites must agree on a common order of serialization.
 - This can be achieved by
 - ✓ Assigning timestamps to transactions
 - ✓ Then having every site serialize conflicting transactions by their timestamps.

Lock Based Algorithms

- In lock based concurrency control algorithms, a transaction must lock a data object before accessing it ,
- In a locking environment, a transaction T is a sequence $\{a_1(d_1), a_2(d_2), \dots, a_n(d_n)\}$ of n actions,
 - ✓ a_i is the operation performed in the i^{th} action
 - ✓ the d_i is the data object acted upon in i^{th} action.
- In addition to read and write, lock and unlock are also permissible actions in locking algorithms.

Lock Based Algorithms

- . A transaction can lock a data object d_i with a “lock(d_i)” action and can relinquish the lock on d_i by an “unlock(d_i)” action.
- A log that results from an execution where a transaction attempting to lock an already locked data object waits, is referred to as a legal log.

Lock Based Algorithms

- . A transaction is well-formed if it
 - ✓ Locks a data object before accessing it,
 - ✓ Does not lock a data object more than once, and
 - ✓ Unlocks all the locked data objects before it completes.

Static Locking

- In static locking, a transaction acquires locks on all the data objects it needs before executing any action on the data objects.
- Static locking requires a transaction to predeclare all the data objects it needs for execution.
- A transaction unlocks all the locked data objects only after it has executed all of its actions.
- Static locking is conceptually very simple.

Static Locking

- However, it seriously limits concurrency because any two transactions that have a conflict must execute serially.
- This may significantly limit the performance of the underlying database system.
- Another drawback of static locking is that it requires a priori knowledge of the data objects to be accessed by transactions.
- This may be impractical in applications where the next data object to be locked depends upon the value of another data object.

Basic Timestamp ordering algorithm

- In the basic timestamp ordering algorithm (BTO), the scheduler at each DM keeps track of the largest timestamp of any read and write processed thus far for each data object.
- Let us denote these timestamps by $R\text{-ts}(\text{object})$ and $W\text{-ts}(\text{object})$, respectively.
- Let $\text{read}(x, TS)$ and $\text{write}(x, v, TS)$ denote a read and a write request with timestamp TS on a data object x .
- A $\text{read}(x, TS)$ request is handled in the following manner:
- If $TS < W\text{-ts}(x)$, then the read request is rejected and the corresponding transaction is aborted, otherwise it is executed and $R\text{-ts}(x)$ is set to $\max\{R\text{-ts}(x), TS\}$.

-
- A write(x, v, TS) request is handled in the following manner: If $TS < R\text{-ts}(x)$ or $TS < W\text{-ts}(x)$, then the write request is rejected, otherwise it is executed and $W\text{-ts}(x)$ is set to TS .
 - If a transaction is aborted, it is restarted with a new timestamp.
 - This method of restart can result in a cyclic restart where a transaction can repeatedly restart and abort without ever completing.
 - This algorithm has storage overhead for maintaining timestamps

Thomas Write Rule (TWR)

- The Thomas write rule (TWR) is suitable only for the execution of write actions
- For a $\text{write}(x, v, TS)$, if $TS < W\text{-ts}(x)$, then TWR says that instead of rejecting the write, simply ignore it.
- This is sufficient to enforce synchronization among writes because the effect of ignoring an obsolete write request is the same as executing all writes in their timestamp order.
- However, an additional mechanism is needed for synchronization between reads and writes because TWR takes care of only write-write synchronization.
- Note that TWR is an improvement over the BTO algorithm because it reduces the number of transaction aborts.

1. Whenever a Transaction T issues a **W_item(X)** operation, check the following conditions:

- If $R_TS(X) > TS(T)$ or if $W_TS(X) > TS(T)$, then abort and rollback T and reject the operation. **else**,
- Execute $W_item(X)$ operation of T and set $W_TS(X)$ to $TS(T)$.

2. Whenever a Transaction T issues a **R_item(X)** operation, check the following conditions:

- If $W_TS(X) > TS(T)$, then abort and reject T and reject the operation, else

If $W_TS(X) \leq TS(T)$, then execute the $R_item(X)$ operation of T and set $R_TS(X)$ to the larger of $TS(T)$ and current $R_TS(X)$.

• T1(10) T2(20) T3(30)

• R1(A)

• W1(A)

•

R2(A)

•

W3(A)

• W1(A)

Multiversion Timestamp ordering algorithm

- In the multiversion timestamp ordering (MTO) algorithm,
- A set of R-ts's and $\langle \text{W-ts}, \text{value} \rangle$ pairs (called versions) is kept for each data object at the respective DM's.
- The R-ts's of a data object keep track of the timestamps of all the executed read operations, and the versions keep track of the timestamp and the value of all the executed write operations.
- Read and write actions are executed in the following manner:
 - A read(x , TS) request is executed by reading the version of x with the largest timestamp less than TS and adding TS to the x 's set of R-ts's. A read request is never rejected.

-
- A write(x, v, TS) request is executed in the following way:
 - If there exists a $R\text{-ts}(x)$ in the interval from TS to the smallest $W\text{-ts}(x)$ that is larger than TS , then the write is rejected,
 - otherwise it is accepted and a new version of x is created with time- stamp TS .

Conservative timestamp ordering algorithm

- The conservative timestamp ordering algorithm (CTO) altogether eliminates aborts and restarts of transactions by executing the requests in strict timestamp order at all DM's.
- A scheduler processes a request when it is sure that there is no other request with a smaller (older) timestamp in the system.
- Each scheduler maintains two queues—a R-queue and a W-queue—per TM.
- These queues, respectively, hold read and write requests.
- A TM sends requests to schedulers in timestamp order and the communication medium is order preserving.
- A scheduler puts a new read or write request in the corresponding queue in timestamp order.

➤ This algorithm executes read and write actions in the following way:

✓ A $\text{read}(x, TS)$ request is executed in the following way.

If every W-queue is nonempty and the first write on each W-queue has a timestamp greater than TS, then the read is executed, otherwise the $\text{read}(x, TS)$ request is buffered in the R-queue.

✓ A $\text{write}(x, v, TS)$ request with timestamp TS is executed in the following manner.

✓ If all R-queues and all W-queues are nonempty and the first read on each R-queue has a timestamp greater than TS and the first write on each W-queue has a timestamp greater than TS, then the write is executed, otherwise the $\text{write}(x, v, TS)$ request is buffered in the W-queue.

✓ When any read or write request is buffered or executed, buffered requests are tested to see if any of them can be executed.

✓ That is, if any of the requests in R-queue or W-queue satisfies condition 1 or 2.